

Deadlock Freedom for Session Types Using Separation Logic

Jules Jacobs

Radboud University Nijmegen
julesjacobs@gmail.com

Overview

- ▶ Channels for concurrency
- ▶ Operational semantics
- ▶ Session types
- ▶ Typing channel references
- ▶ Typing channel references with separation logic
- ▶ The connectivity tree
- ▶ Deadlock freedom

Channels for concurrency

- ▶ Lambda calculus + channels
- ▶ Fork: `let c' = fork(fun c => ...)`
- ▶ Send: `let c' = send(c, 3)`
- ▶ Receive: `let (c', x) = recv(c)`
- ▶ Close: `let () = close(c)`

Channels for concurrency

Example:

```
let c' = fork(fun c =>
    send(c, 3)
    let k = recv(c)
    if k < 10 then send(c, "hello") else send(c, "hi")
    close(c)
)

let n = recv(c')
send(c', 2*n)
let msg = recv(c')
close(c')
print(msg)
```

Potential bugs

Safety bugs:

- ▶ Wrong type of message, e.g. string instead of int
- ▶ Example:
- ▶ Using a channel after close
- ▶ Example:

Concurrency bugs:

- ▶ Not sending a message (receive blocks \implies deadlock)
- ▶ Both sides do a receive
- ▶ Example:

Leaks:

- ▶ Not closing a channel (leak)
- ▶ Example:

Session types

Existing system, citation: Honda et al. 1900

Session types

- ▶ Linear lambda calculus
- ▶ Channels have protocol types: sequence of send (!) receive (?) and End
- ▶ Example: !Int, ?Int, !String, End
- ▶ Dual: ?Int, !Int, ?String, End
- ▶ If $c : !T, R$ and $x : T$ then **send**(c,x) : R
- ▶ If $c : ?T, R$ then **recv**(c,x) : R x T

State passing

Example:

```
let c' = fork(fun c =>
  let c = send(c, 3)
  let (c,k) = recv(c)
  if k < 10 then
    let c = send(c,"hello")
    close(c)
  else
    let c = send(c,"hi")
    close(c)
)
let (c',n) = recv(c')
let c' = send(c', 2*n)
let (c',msg) = recv(c')
close(c')
print(msg)
```


Session types prevent bugs

Safety bugs:

- ▶ Wrong type of message, e.g. string instead of int
- ▶ Example:
- ▶ Using a channel after close
- ▶ Example:

Concurrency bugs:

- ▶ Not sending a message (receive blocks \implies deadlock)
- ▶ Both sides do a receive
- ▶ Example:

Leaks:

- ▶ Not closing a channel (leak)
- ▶ Example:

Problem statement

- ▶ Deadlock freedom for session types has not been mechanised
- ▶ Lambda calculus
- ▶ Cut elimination

Key ideas:

- ▶ Research accounting via separation logic
- ▶ Connectivity graphs for deadlocks

Operational semantics

- ▶ Small-step operational semantics
- ▶ State: threads \times buffers
- ▶ Buffers indexed by addresses (e.g. natural numbers); two buffers per channel
- ▶ Representation of a channel is address \times bool, e.g. `send(#(324, true))`
- ▶ `send(#(324, true), "hello")` \mapsto `send(#(324, true)`
- ▶ Send adds message to the other party's buffer
- ▶ Receive takes a message out of its own buffer
- ▶ Fork allocates a new pair of buffers, returns `send(#(n, true))` to one side and passes `send(#(n, false))` to the other
- ▶ Close deallocates its own buffer

```
endpoint = nat x bool
heap = endpoint -> list val
local_step : expr x heap -> expr x heap -> Prop
step : list expr x heap -> list expr x heap -> Prop
```

Formal properties

Formal properties:

- ▶ Safety
- ▶ Deadlock freedom
- ▶ Leak freedom

Theorem

Typed expr \vdash previous 3 properties

Progress & preservation

Formal properties

High level approach:

- ▶ Progress & preservation (explain this) \implies safety
- ▶ Progress & preservation for linear types
- ▶ Connectivity tree

Typing channel references


- ▶ When the operational semantics takes steps, channel references (e.g. $\#(324, \text{true})$) enter into the program.
- ▶ How do we type them?
- ▶ We need a typing rule for channel references.
- ▶ We need an environment Σ for channel references, similar to Γ for variables
- ▶ If $((342, \text{true}), T) \in \Sigma$ then $\#(324, \text{true}) : T$

Typing channel references with separation logic

How to explain this...?

The connectivity graph

- ▶ A graph containing threads (boxes) and channels (circles)
- ▶ An arrow $T \rightarrow C$ means that thread T has a reference to channel C
 - ▶ This means that some data structure that T holds contains a channel reference to C
- ▶ An arrow $C1 \rightarrow C2$ means that channel $C1$ has a reference to channel $C2$
 - ▶ This means that somewhere in the buffer of $C1$ there is a reference to $C2$



Deadlock freedom

- ▶ Invariant: the connectivity graph is a tree.
- ▶ Even when channel references are put in buffers.
- ▶ Ensured by linearity.

Deadlock freedom

Theorem

There is a thread that's not blocked.

Proof.

We find a thread that's not blocked using this algorithm:

- ▶ Start at any thread in the connectivity graph.
- ▶ If the thread is not blocked, we're done.
- ▶ If the thread is blocked, step to the channel that it's blocked on.
- ▶ If we're at a channel

